

Trust and AI safety

Week 3.

NOTES · 1 / 8

Session of Monday 21 September. AIBS theme: trust and AI safety. AEL decision: where the team's knowledge lives.

Desk research gets a tool; field research gets new data

DESK RESEARCH

A new way of working: you design the tool that helps you.

FIELD RESEARCH

New data, external validation, contact with your target audience.

NOTES · 2 / 8

Desk research: the tool AEL is building for the team's own research is the new part of the "way of working" — not a shortcut, a designed instrument. Field research: this is the case study visit, and it earns its place by giving the team

something a desk cannot — data nobody else has, and a check against everything the desk work assumed.

How do you create value in desk and field research?

DESK RESEARCH

- **Reliable?** Same result if you run it again.
- **Valid?** Measures what it claims to.
- **Transparent?** You can show how it works.

FIELD RESEARCH

Relevance? What your case study says beyond itself.

NOTES · 3 / 8

Reliability, validity and transparency are the three questions to ask of a tool's own output, not just of what it retrieves. A team that cannot say "yes" to all three does not yet have a research instrument, it has a demo. Field research earns its

place on a different axis: not whether it is reliable in the same sense, but whether the one case generalises — its relevance beyond itself.

Reliability, validity, transparency — your quality criteria

What changes when you work with AI?

NOTES · 4 / 8

The pivot slide. Reliability, validity and transparency are not new to research — they are the field's existing quality bar. What is new is asking them of a process

that includes an AI tool: the tool's output has to clear the same bar a human researcher's would, and a team should be able to say how it knows.

Discussion: your first handbook page

- **Search strategy.** How did you use it in the tool?
- **Quality criteria.** Were they respected?
- **Output quality.** How would you rate it today?

NOTES · 5 / 8

The room's own first result, not a hypothetical: this is the first time the search-strategy and quality-criteria assignments from week 2 come back to be checked rather than just handed in. Push on the middle question — "did you check" is a

different claim from "did you write criteria" — and let the third question be answered honestly rather than diplomatically; a flat "not yet" is a better answer here than a confident guess.

Trust splits three ways: the tool, you, and ownership

- **The tool.** Can you trust what it produces?
- **You.** Can you be trusted as the researcher?
- **Ownership.** Whose output is it?

NOTES · 6 / 8

Three separate questions, deliberately not collapsed into one. Trusting the tool's output is a methods question. Being trusted as a researcher is about the team's own credibility once AI is part of the process — a reader who cannot tell what the

team checked will trust the whole page less, not more. Ownership is the one students usually have not considered: if the tool drafted it, whose work is it, and what do you say about that in the write-up?

One exercise on your own, before week 4

About 25 minutes, any evening this week. The exercise is here.

- **Bring** your own AI assistant, signed in.
- **Your team number** sets the order of two rounds. Nothing else.
- **In week 4** we read what the whole cohort did.

NOTES · 7 / 8

The exercise students do at home this week, alone, before the week 4 session. Say what it costs them (about 25 minutes and a laptop) and that the team number only sets the order of the two rounds. Say nothing about what the two

rounds differ in: the comparison only works if they meet it cold. Their results come back in week 4, as one picture of the whole cohort.

The rules we keep, week by week

Every time we spot a habit worth keeping, we turn it into a rule.

RULE 1

Before taking a break, **always** set AI to work.

RULE 2

Do something **weird**.

RULE 3

Added together in the session.

NOTES · 8 / 8

The rulebook, not the recap: read Rule 1 and Rule 2 as fixed and Rule 3 as the one the room writes today, the way week 2 wrote Rule 2. If nobody proposes

one, offer a candidate rather than close the slide empty.