

Trust and AI safety

Week 3.

NOTES · 1 / 11

Session of Monday 21 September. AIBS theme: trust and AI safety. AEL decision: where the team's knowledge lives.

The experiment, debriefed

Added with the rest of the week 3 deck.

NOTES · 2 / 11

Placeholder for the experiment debrief. It stays empty until the session itself supplies the result — the debrief is read off the room's own data, not written in advance. The AI agents segment (s8–s10) still comes after the experiment rather than before it: an agents segment ahead of the experiment would be a treatment

nobody randomised. The research-methodology segment (s3–s7) sits between the two safely, because it runs after the experiment has already been played, not before.

Desk research gets a tool; field research gets new data

DESK RESEARCH

A new way of working: you design the tool that helps you.

FIELD RESEARCH

New data, external validation, contact with your target audience.

NOTES · 3 / 11

Desk research: the tool AEL is building for the team's own research is the new part of the "way of working" — not a shortcut, a designed instrument. Field research: this is the case study visit, and it earns its place by giving the team

something a desk cannot — data nobody else has, and a check against everything the desk work assumed.

How do you create value in desk and field research?

DESK RESEARCH

- **Reliable?** Same result if you run it again.
- **Valid?** Measures what it claims to.
- **Transparent?** You can show how it works.

FIELD RESEARCH

Relevance? What your case study says beyond itself.

NOTES · 4 / 11

Reliability, validity and transparency are the three questions to ask of a tool's own output, not just of what it retrieves. A team that cannot say "yes" to all three does not yet have a research instrument, it has a demo. Field research earns its

place on a different axis: not whether it is reliable in the same sense, but whether the one case generalises — its relevance beyond itself.

Reliability, validity, transparency — your quality criteria

What changes when you work with AI?

NOTES · 5 / 11

The pivot slide. Reliability, validity and transparency are not new to research — they are the field's existing quality bar. What is new is asking them of a process

that includes an AI tool: the tool's output has to clear the same bar a human researcher's would, and a team should be able to say how it knows.

Discussion: your first handbook page

- **Search strategy.** How did you use it in the tool?
- **Quality criteria.** Were they respected?
- **Output quality.** How would you rate it today?

NOTES · 6 / 11

The room's own first result, not a hypothetical: this is the first time the search-strategy and quality-criteria assignments from week 2 come back to be checked rather than just handed in. Push on the middle question — "did you check" is a

different claim from "did you write criteria" — and let the third question be answered honestly rather than diplomatically; a flat "not yet" is a better answer here than a confident guess.

Trust splits three ways: the tool, you, and ownership

- **The tool.** Can you trust what it produces?
- **You.** Can you be trusted as the researcher?
- **Ownership.** Whose output is it?

NOTES · 7 / 11

Three separate questions, deliberately not collapsed into one. Trusting the tool's output is a methods question. Being trusted as a researcher is about the team's own credibility once AI is part of the process — a reader who cannot tell what the

team checked will trust the whole page less, not more. Ownership is the one students usually have not considered: if the tool drafted it, whose work is it, and what do you say about that in the write-up?

You have used an agent since week 1

Agent: an LLM (large language model) using tools in a loop, towards a goal.

- **Workflow or agent?** Code or model: who picks the next step?
- **The loop:** reason, act, observe (ReAct).
- **The week 1 intern:** capable, needs oversight.

NOTES · 8 / 11

About five minutes. Connects to AEL week 1 (the spike) and AEL week 2 (the debrief that named ReAct).

Every team has used an agent for two weeks: the agentic command line of the week 1 spike. This card gives that a name. The definition is the one Google and Anthropic developer material both use by mid-2026: an LLM with tools, running in a loop, towards a goal. Wooldridge's textbook definition from 2009 leaves the LLM out. A thermostat passes his test.

The line between a workflow and an agent is who picks the next step. In a workflow, your code sets the path. In an agent, the model decides. Both can use tools and loop.

The loop is the one they watched in the spike: reason, act, observe, and again. If the week 2 debrief did not name ReAct, introduce it here rather than recall it.

The intern from week 1 still holds: capable, and still needs someone in charge of judgement and oversight — Karpathy's phrase for it is "intern entities".

The model is rented; the harness is the part you own

- **Harness:** the software around the model. Model as CPU, harness as operating system.
- **This week's choice:** where knowledge lives. RAG retrieves, LLM Wiki compiles, Fat Skills act — one question, four ways.
- **Still ahead:** Contracts, Constraints, Compounding.

NOTES · 9 / 11

About five minutes. Connects to AEL week 3, this week's knowledge-architecture decision, and back to AIBS week 2 on vendor dependency.

The harness is everything around the model: what it sees, which tools it may call, what happens after each step. Open the *Harness is the OS* visualisation: the model is the CPU, "powerful but inert" in the source's words; the context window is the RAM; the harness is the operating system that decides what the CPU sees and when.

Why it matters to an SME: the model is rented from a vendor who will be overtaken within the year. The harness is what a company builds, owns and

improves.

This week's AEL decision is one part of that harness — where the team's knowledge lives. RAG retrieves passages from documents at the moment of the question. An LLM Wiki compiles sources into pages ahead of time. Fat Skills act on what is known without being asked. Open *one question, four ways* and let the room see the same question answered each way. The choice follows from what the agent's job is.

Name Contracts, Constraints and Compounding and move on. They come in weeks 4 to 6, each after the failure it answers.

Trusting an agent splits three ways: data, output, people

- **Data.** Rent the model, or run a small one in-house.
- **Output.** Users regret an agent that acted beyond what they authorised — even when it was right.
- **People.** Hiding AI use: 14 % where trust is highest, 47 % where lowest.

NOTES · 10 / 11

About five minutes. Connects to AIBS week 3: the brief asks whether the SME can trust AI, and splits the answer three ways.

Data. The model is rented and the harness owned, so the question is what leaves the building. A small language model is defined by the device it fits on, which turns "can we run it ourselves?" into a question with an answer.

Output. People regret an agent that acted beyond what they would have authorised, even when the result was correct. The trigger is an action that cannot be recalled and that someone outside sees — sending an email is the

example. Tie it to the room: in one of the two rounds everyone put their position to an assistant. By the definition on the first agents card that was not an agent — no tools, no loop, it could only reply. Nothing it did could be recalled or seen outside the room, which is the safe end of that scale. Where should an SME's first real agent sit?

People. Whether employees admit to using AI follows trust. In a survey of 604 daily users, 14 % in the most-trusting quartile hid their AI use, against 47 % in the least-trusting.

The rules we keep, week by week

Every time we spot a habit worth keeping, we turn it into a rule.

RULE 1

Before taking a break, **always** set AI to work.

RULE 2

Do something **weird**.

RULE 3

Added together in the session.

NOTES · 11 / 11

The rulebook, not the recap: read Rule 1 and Rule 2 as fixed and Rule 3 as the one the room writes today, the way week 2 wrote Rule 2. If nobody proposes

one, offer a candidate rather than close the slide empty.