

AI use cases at SMEs

Week 5.

NOTES · 1 / 6

Session of Monday 5 October. AIBS theme: AI use cases at SMEs. AEL: one way to keep knowledge, built live from an empty folder.

What small firms really use AI for

The AIBS segment is added here before Monday.

An LLM wiki reads each source once and keeps what it read

RAW/

The sources. Never edited.

WIKI/

The pages. The assistant writes.

AGENTS.MD

The rules. You change them.

Karpathy, LLM Wiki, April 2026.

NOTES · 3 / 6

Most of what the teams built in week 4 looks things up when a question is asked: files uploaded to a chat, a folder the assistant searches, a vector store. That works, and it forgets. Ask twice and it does the work twice; a contradiction between two sources is found only if a question happens to touch both. Karpathy: the LLM "is rediscovering knowledge from scratch on every question". The week 3 visualisation (one question, four ways) is the reminder that this is one of four options.

Then the three folders, three owners. raw/ is the evidence and nobody edits it, the assistant included. wiki/ belongs to the assistant; the team reads and asks for changes. AGENTS.md is the contract, and the team changes it. All six assistants the students use read AGENTS.md; CLAUDE.md and GEMINI.md point to it.

Three sources, three formats, one empty wiki

- **Article.** [OECD Cogito blog](#), 2025.
- **Video.** [Asian Productivity Organization](#), 64 minutes.
- **Report.** [OECD survey](#), 2026, 36 pages.

NOTES · 4 / 6

The demo. Leave this card up and switch to the terminal and Obsidian.

Before class (the day before). Create the demo repository from the template, install the helpers, and acquire and ingest the video and the report. They are too long to run live: in the dry run on 30 September (headless, no pause) the video took 10½ minutes and the report 17. Commit after each. The article stays in `raw/` un-ingested, or is fetched live.

Acquire, one command per format (show at least the article live; the other two take seconds):

- Article: ``python3 tools/fetch_article.py https://oecdcoigito.blog/2025/09/16/agentiic-ai-for-small-business-growth/`` — about 1,800 words, word for word.
- Video: ``python3 tools/fetch_youtube.py https://www.youtube.com/watch?v=51IXx4wBuHE`` — captions only; 64 minutes, about 7,600 words of automatic captions.
- Report: save the PDF in `raw/reports/`, then ``markitdown raw/reports/oecd-empowering-smes-in-the-age-of-ai.pdf > raw/reports/oecd-empowering-smes-in-the-age-of-ai.md`` — 36 pages, about 13,000 words, four seconds. The PDF itself is not committed (.gitignore).

Live: the article. ``/ingest raw/articles/agentiic-ai-for-small-business-growth.md`` — added by Witek`. Let the pause happen and read the takeaways aloud before saying go. On an empty wiki it took just under five minutes; on top of the report and

the video expect longer, so budget ten with the pause and talk. Because the survey is already in, the disagreement appears live.

Four things to show:

1. The pause: the place to catch a misreading before it lands on ten pages. 2. The source page's caveats. In the dry run it found that the author heads PayPal's government relations, that PayPal is a partner of the survey the article cites, that the "72%" figure links back to the article itself, and that the two business owners are invented examples. 3. The disagreement. In the dry run it landed on the `agentiic-ai` concept page under Debates: the article writes in the present tense about agents that reorder stock and negotiate with suppliers; the OECD survey counts 3.6% of AI users running agentiic AI, in a sample it says is not representative and skews to the digitally mature. The page says what the disagreement turns on (partly tense, partly being found by other people's agents versus running your own) and the page's confidence dropped to 0.65. 4. Query and lint. Ask: `*Which AI uses pay off first for a small manufacturer, according to the wiki?*` Show the citations, then the query entry in the log with the pages it read. Then ``/lint``, and dwell on the quote check: every quotation on a wiki page is compared with the raw file, so an invented quotation fails. These are the two habits the knowledge architecture criteria said would come back in week 5.

Fallback. The dry-run wiki, all three sources ingested and lint clean (52 pages), is on the lecturer's laptop outside the repositories. Open it in Obsidian if the network or the model fails.

Your team's copy takes ten minutes

[datadrivendecisions/llm-wiki-template](#)

Never interview material.

NOTES · 5 / 6

Teams make their copy now if they did not before the session; the README has the steps. Walk the room: the usual problems are Python on Windows (‘python’, not ‘python3’) and a venv that is not activated. Each team ingests one published source of its own and reads the source page back. Never interview material, private repository or not: what the assistant reads goes to a model service.

Nobody has to switch. A folder with a naming rule, traced end to end, is still a good answer. What the wiki costs: every source is read when it arrives, whether anyone asks or not; a misreading lands on several pages at once; past a few hundred pages the index is not enough and you need search. A team that adopts it writes a decision-log entry and versions its knowledge architecture.

Before 12 October: the first evaluation, with your decision log

The assignment is here.

NOTES · 6 / 6

The first evaluation is held against the PRD's own criteria, and the decision log travels with it. The lookup record and the quotation check are what let a team explain a wrong sentence on its page instead of guessing.