

AI use cases at SMEs

Week 5.

NOTES · 1 / 11

Session of Monday 5 October. AIBS theme: AI use cases at SMEs. AEL: one way to keep knowledge, built live from an empty folder.

What small firms really use AI for

The AIBS segment is added here before Monday.

Most tools read your sources again for every question

"The LLM is rediscovering knowledge from scratch on every question."

Andrej Karpathy, LLM Wiki, April 2026. Week 3: one question, four ways.

NOTES · 3 / 11

Most of what the teams built in week 4 looks things up at the moment a question is asked: files uploaded to a chat, a folder the assistant searches, a vector store. That works, and it forgets. Ask the same question twice and it does the work

twice; a contradiction between two sources is found only if a question happens to touch both. Karpathy's point is that nothing accumulates. The week 3 visualisation is the reminder that this is one of four options, not the new default.

An LLM wiki reads each source once and keeps what it read

RAW/

The sources. You add them; nobody edits them.

WIKI/

The pages. The assistant writes; you check.

AGENTS.MD

The rules. You change them.

NOTES · 4 / 11

Three folders, three owners. raw/ is the evidence and is never edited, the assistant included. wiki/ belongs to the assistant: the team reads it and asks for changes. AGENTS.md is the contract, and the team changes it when it does not

fit. All six assistants the students use read AGENTS.md; CLAUDE.md and GEMINI.md in the template point to it.

Three sources, three formats, one empty wiki

- **Article.** OECD Cogito blog, 2025.
- **Video.** Asian Productivity Organization, 64 minutes.
- **Report.** OECD survey, 2026, 36 pages.

NOTES · 5 / 11

Demo, part 1 (acquire). One command per format, each lands a markdown file in raw/:

- Article: ``python3 tools/fetch_article.py https://oecdcoigito.blog/2025/09/16/agentice-ai-for-small-business-growth/`` — about 1,800 words, verbatim.
- Video: ``python3 tools/fetch_youtube.py https://www.youtube.com/watch?v=51IXx4wBuHE`` — captions only, the video is not downloaded; 64 minutes,

about 7,600 words of automatic captions.

- Report: save the PDF in raw/reports/, then ``markdown raw/reports/oecd-empowering-smes-in-the-age-of-ai.pdf > raw/reports/oecd-empowering-smes-in-the-age-of-ai.md`` — 36 pages, about 13,000 words, four seconds.

Point out that the PDF itself is not committed (.gitignore): the text is, and the source page says where the original lives.

Watch the pause: it reads, then asks before it writes

- **The pause.** Your chance to catch a misreading.
- **The source page.** Who wrote it; who paid for it.
- **Index and log.** Every page, every step.

NOTES · 6 / 11

Demo, part 2 (process). In Claude Code: ``/ingest raw/articles/agent-ai-for-small-business-growth.md — added by Witek``. Let the pause happen and read the takeaways out loud before saying go. On the source page, show the byline: a PayPal author on the OECD's blog, and the two business owners in it are invented examples. Then open `wiki/index.md` and `wiki/log.md`.

Dry run on 30 September, headless, no pause: the article took just under five minutes and wrote twelve pages. With the pause and talking, budget eight. Ingest the article live; if time is short, run the video and the report while talking over s7.

Two sources, two pictures of AI at small firms

THE ARTICLE

An AI assistant that “negotiates with suppliers”.

Written by someone from PayPal.

THE REPORT

61 % of SMEs use AI. **76 %** of those are “AI novices”.

OECD survey, 2,000+ SMEs.

NOTES · 7 / 11

The disagreement the wiki should surface once both are in: the article describes an assistant that runs inventory and negotiates with suppliers; the OECD's own survey finds most SME users are novices using off-the-shelf tools. Show where

it lands — the Debates section of the concept page, and the confidence number that went down. Also worth a sentence: the OECD calls its own sample non-representative, and the source page should say so.

Every answer shows its sources, and the log shows what it read

- **Query.** A link after every claim.
- **Log.** The pages each answer read.
- **Lint.** Every quotation checked against its source.

NOTES · 8 / 11

Demo, part 3 (query and lint). Ask: *Which AI uses pay off first for a small manufacturer, according to the wiki?* Show the citations, then the query entry in the log with the pages it read. Then `/lint`. The quote check is the one to dwell

on: it compares every quotation on a wiki page with the raw file, so an invented quotation fails. These are the two habits the knowledge architecture criteria said would come back in week 5.

Your team's copy takes ten minutes

[datadrivendecisions/llm-wiki-template](#)

- **Use this template** → a new repository.
- **Clone, install,** run the check.
- **Ingest** one published source.

Never interview material.

NOTES · 9 / 11

Teams make their copy now, if they did not before the session. The README has the steps. Walk the room: the usual problems are Python not on the PATH on

Windows (`python` instead of `python3`) and the venv not activated. Each team ingests one published source of its own and reads the source page back.

Adopting the wiki is a decision, and it has costs

- **Reading up front.** Every source costs model time.
- **A misreading spreads.** Several pages at once.
- **A few hundred pages.** Then you need search.

Adopting it? An entry in your decision log.

NOTES · 10 / 11

Nobody has to switch. A folder with a naming rule, traced end to end, is still a good answer. What the wiki costs: every source is read when it arrives whether anyone asks or not; a misreading lands on several pages at once; past a few

hundred pages the index is not enough. A team that adopts it writes a decision-log entry and versions its knowledge architecture.

Before 12 October: the first evaluation, with your decision log

The assignment is here.

- **Held against** your PRD's criteria.
- **A record** of each lookup.
- **Every quotation** checked against its source.

NOTES · 11 / 11

The first evaluation is held against the PRD's own criteria. The decision log travels with it. The lookup record and the quotation check are what let a team

explain a wrong sentence on its page instead of guessing.